

UCSD CSE131 F19 – Diamondback

October 24, 2019

Due Date: 11pm Wednesday, November 6

Closed to Collaboration

You will implement Diamondback, a language with functions and a static type system.

Classroom: <https://classroom.github.com/a/TyAjQxca> Github: <https://github.com/ucsd-cse131-f19/pa4-student>

Syntax

The concrete syntax and type language for Diamondback is below. We use \dots to indicate *zero or more* of the previous element¹. So a program p is a sequence of zero or more definitions d followed by an expression e , and a definition has zero or more arguments $x:\tau$ and one or more expressions in its body $e\dots$. There are boxes around the new pieces of concrete syntax.

e	$:=$	$n \mid \text{true} \mid \text{false} \mid x$ $\mid (\text{let } ((x\ e)\ (x\ e)\ \dots)\ e\ e\dots)$ $\mid (\text{if } e\ e\ e)$ $\mid (op_2\ e\ e) \mid (op_1\ e)$ $\mid (\text{while } e\ e\ e\dots) \mid (\text{set } x\ e)$ $\mid \boxed{(f\ e\dots)}$		
d	$:=$	$\boxed{(\text{def } f\ (x:\tau\dots):\tau\ e\ e\dots)}$	τ	$:=$ Num Bool
p	$:=$	$\boxed{d\dots e}$	Δ	$:=$ $\{f:\tau\dots \rightarrow \tau, \dots\}$
op_1	$:=$	add1 sub1 isNum isBool print	$\Delta[f]$	means look up the type of f in Δ
op_2	$:=$	+ - * < > ==	Γ	$:=$ $\{x:\tau, \dots\}$
n	$:=$	63-bit signed number literals	$\Gamma[x]$	means look up the type of x in Γ
x, f	$:=$	variable and function names	$(x, \tau) :: \Gamma$	means add x to Γ with type τ
			$\Delta; \Gamma \vdash e : \tau$	means with definitions Δ and env Γ , e has type τ
			$\Delta \vdash_d d : \checkmark$	means with definitions Δ the definition d type-checks
			$\vdash_p p : \checkmark$	means the program p type checks

There are two different **def** and function application forms because functions are allowed to have zero arguments, but the \dots notation means one or more. You can easily represent arguments in both cases with a list that's allowed to be empty.

An example program that computes whether a number is even or odd extremely inefficiently and prints several examples is:

```
(def even (n : Num) : Bool
  (if (== n 0) true (odd (- n 1))))
(def odd (n : Num) : Bool
  (if (== n 0) false (even (- n 1))))
(def test() : Bool
  (print (even 30))
  (print (odd 30))
  (print (even 57))
  (print (odd 57)))

(test)
```

Expected output:
true
false
false
true
true

¹This is different than the last assignment, but important for avoiding some notational clutter.

Semantics

Function Definitions and Applications

The main new feature in Diamondback is function definitions d and function applications $(f\ e\cdots)$. A function application $(f\ e_a\cdots)$ uses the function definition with the matching name (`def $f\ (x\ :\tau\cdots)\ e\cdots$`), and should evaluate to the same result as `(let (($x\ e_b$) \cdots) $e\cdots$)`, where the argument expressions $e_a\cdots$ come from the application, the names $x\cdots$ come from the definition, and the body expressions $e_b\cdots$ come from the definition.²³

The Main Expression

Diamondback programs are expected to have a single expression after the definition list that is the main entry point for the program. This expression has `input` bound to the user’s input by default.

As an example, consider this program:

```
(def (abs x : Num) : Num
  (if (< x 0) (* -1 x) x))
(* (abs input) 2)
```

We could think of its evaluation taking these steps (with an input of 5):

```
→ (* (abs input) 2)
→ (* (abs 5) 2)
→ (* (let ((x 5)) (if (< x 0) (* -1 x) x)) 2)
→ (* (if (< 5 0) (* -1 5) 5) 2)
→ (* (if false (* -1 5) 5) 2)
→ (* 5 2)
→ 10
```

Printing

Diamondback also adds a new primitive, `print`, that prints a value to the console followed by a newline. Numbers should print as their user-interpreted value, so `(print 22)` should print 22, and `(print true)` should print `true`.

The entire `print` expression evaluates to the same value as its argument (which is the value that gets printed), so `(print (+ 1 7))` evaluates to 8 in addition to printing it.

Type Checking

Diamondback has essentially the same type rules as Copperhead for expressions, with two changes. First, all of the rules contain a definitions environment Δ in addition to the type environment Γ . Second, there are two new rules:

$$\text{TR-PRINT} \frac{\Gamma \vdash e : \tau}{\Delta; \Gamma \vdash (\text{print } e) : \tau} \quad \text{TR-APP} \frac{\Delta[f] = \tau_1 \cdots \tau_n \rightarrow \tau_r \quad (\Delta; \Gamma \vdash e_1 : \tau_1) \cdots (\Delta; \Gamma \vdash e_n : \tau_n)}{\Delta; \Gamma \vdash (f\ e_1 \cdots e_n) : \tau_r}$$

²³Note that replacing application expressions with let expressions is *not* a strategy that works in general in the compiler, because the body expressions $e_b\cdots$ could contain other function applications. This description is, however, a useful way to describe the behavior of a function application succinctly and is a perfectly valid way to evaluate functions “by hand” when we can write out all the intermediate steps with concrete values. See the reading for more detail: <https://ucsd-cse131-f19.github.io/lectures/10-22-lec8/notes.pdf>

³We actually need to be a little bit careful here. While this rule worked fine for the single-argument functions in the reading, here we’d really want a version of let that doesn’t include earlier bindings in later ones to avoid clashes between names in scope and names in the function’s argument list. This is really only relevant if you’re writing things out on paper, and doesn’t affect the design of a calling convention at all.

TR-Print says that **print** expressions' result is the same type τ as the argument, which matches the semantics.

In English, TR-App rule says

If the function definition f has argument types $\tau_1 \cdots$ and return type τ_r , and the arguments $e_1 \cdots$ of an application of f have matching types, then the application has type τ_r .

It's common to write the type of definitions as $\tau_1 \cdots \rightarrow \tau_r$, also called an "arrow type", to be evocative of taking a number of argument types and producing a result type.

The existing rules are unchanged aside from tracking Δ (which only TR-App uses).

There are also two new rules, one for type-checking definitions, and one for type-checking programs. They use slightly different \vdash notation with subscripts $_d$ and $_p$ to indicate that they have meaning for pieces of syntax other than expressions. Since, unlike expressions, we don't calculate their overall type, we simply use \checkmark in the rule to indicate that they pass all checks.

$$\text{TR-DEF} \frac{\Delta; \{x_1 : \tau_1, \dots x_n : \tau_n\} \vdash e \cdots : \tau_r}{\Delta \vdash_d (\text{def } f (x_1 : \tau_1 \cdots x_n : \tau_n) : \tau_r \ e \cdots) : \checkmark}$$

$$\text{TR-PROG} \frac{(\Delta; \vdash_d d_1 : \checkmark) \cdots (\Delta; \vdash_d d_n : \checkmark) \quad \Delta; \{\text{input} : \text{Num}\} \vdash e : \tau}{\vdash_p d_1 \cdots d_n e : \checkmark}$$

Where in \vdash_p , Δ is constructed by mapping each definition name to an arrow type made of its argument types and return type, so $(\text{def } f (x : \tau_1 \cdots x_n : \tau_n) : \tau_r \ e \cdots)$ would appear in Δ as $\{f : \tau_1 \cdots \tau_n \rightarrow \tau_r\}$. In English, TR-Def says that a definition type-checks if its body has the expected return type τ_r when type-checked in an environment with just the arguments of the definition mapped to their declared types. A program p type-checks if all of its definitions type-check and its main expression has some type in the environment that assumes **input** has type **Num** (along with also assuming the declared definitions).

New Errors & Miscellaneous

Type errors should be reported with "Type mismatch" as usual, including type errors resulting from the new rules, including passing the wrong number of arguments to a function. If a function application uses a function name that isn't defined, the compiler should report an error containing "Unbound".

It's allowed for variables and functions to use the same name, so there could be a top-level definition named **f** and a variable in an argument or let named **f**.

It's a well-formedness error for multiple functions to have the same name, or for multiple arguments within the same function to have the same name. Report these cases with an error that contains the string "Multiple functions" and "Multiple bindings" respectively.

An empty function body should be reported with "Invalid" as with other syntax errors.

Implementation Recommendations and Details

Registers Used by main

You may have reasons to want to use registers like **rbx**, **rbp**, **rdi** or others. The assembly generated by gcc and clang for **main** may use these registers as well, so they should be saved at the beginning of **our_code_starts_here** and restored before the final **ret**. You can use **push rbx** and **pop rbx** to accomplish this.

Stack Alignment

Some systems require that the stack pointer `rsp` be aligned at a 16-byte (2 word) boundary before making calls into library functions that use system calls, like `printf`. If you get stack alignment segmentation faults⁴ you may want to make sure your calling convention always moves `rsp` by multiples of 16, which could mean leaving an extra word of space on some calls. For example, if your calling convention uses 2 words for the old value of `rsp` and the return address, then a function call with an *odd* number of arguments could end up on an 8-byte boundary. You can test this by using `print` in functions with varying numbers of arguments.

Moving Labels into Memory

In class, we used code like `mov [rsp-16], after_call` to move a label into memory. This actually requires two instructions on some platforms. If you see an error like "format does not support 32-bit absolute addresses." you may be running into this. The solution is simple, just save the label into a register first:

```
mov rax, after_call
mov [rsp-16], rax
```

Print

While you're free to implement `print` in any way you prefer that works, one that we found expedient is to call a function defined in `main.c`. This requires using C's calling convention. On x86-64, this means moving the argument into register `rdi`, moving `rsp` to free space at the top of the stack, and then using the `call` instruction to push the current code address to the stack and jump to the function you wrote in `main`. On return, `rsp` should be moved back to its original location, and your generated code should ensure that the printed value ends up in `rax`.

Keep in mind that if you use registers other than `rax`, they may be overwritten during the use of `print`. Wikipedia has a reasonable, brief, accurate summary of the callee-save vs. caller-save behavior at https://en.wikipedia.org/wiki/X86_calling_conventions#System_V_AMD64_ABI. We assume that users of Diamondback will use `gcc` and `clang` on systems that use this convention, not the Microsoft convention.

Required Tests

You **must** write the following programs to test your compiler. You can pick any little details you want like function names and base case behavior, but they must be in files with these names:

- `input/fibonacci.boa` – A program that defines a fibonacci function and calls it on `input`.
- `input/remainder.boa` – A program that defines a function that takes two arguments and uses repeated subtraction to get the remainder of the first divided by the second. Test it on several pairs of inputs and include `print` expressions demonstrating it. (Yes, we are asking you to write tests for the function you're writing as a test.)
- `input/isprime.boa` – A program that uses the definition from `remainder.boa` (you can copy/paste it) along with a loop to test if `input` is prime, and print `true` if it is and `false` if not. Your algorithm does not have to be sophisticated; it's fine to try the numbers from 2 to `input` and check the remainder.
- `input/deepstack.boa` – A program of your choice that has at least 3 different functions, each with a different number of arguments. When run, it should require a runtime stack of at least depth 10, where each function has at least one frame represented on the stack.

These tests together will be worth 20% of your grade (you just have to write them, you still get credit even if they don't fully work).

⁴We saw this in class <https://github.com/ucsd-cse131-f19/ucsd-cse131-f19.github.io/blob/master/lectures/10-10-lec5/compile.ml#L106>

You should start with smaller tests and write many more tests than these. We provide these because they usefully exercise your compiler and demonstrate that you can write some meaningful computations now.

Describing Your Calling Convention

As you implement Diamondback, you will need to make a number of decisions, not least of which is the calling convention you choose, and decisions you make around compiling application expressions and definitions. Along with your code, you will write a design document as a separate PDF describing how your calling convention works.

There are no restrictions on the calling convention you choose. You could implement the x86-64 convention for your functions, the version we discussed in class generalized to arbitrary lists of values, or something of your own design. Different choices will require different implementation strategies, different stack management requirements, and produce code that is “interesting” to debug in various ways.

In your design document, you should make sure to cover (in whatever order makes sense) the items below. This report will be worth **30% of your assignment grade**.

1. A description of your calling convention in general terms:

- (a) What does the caller do before and after the call?
- (b) What is the callee responsible for?
- (c) Are there improvements you can imagine making in the future?

Use snippets of OCaml code from your compiler to illustrate key features.

2. Pick three example programs that use functions and are interesting in different ways (you can use the required tests if you like), and use them to describe your calling convention:
 - (a) Show their source, generated assembly, and output (you can summarize the generated code if it’s quite long)
 - (b) Highlight the parts of the generated assembly that make the example interesting, distinct, and/or especially challenging to compile

Still write this even if you don’t have everything working! In that case, in part 2, pick at least one example that *doesn’t* work, and note both its expected output and its actual behavior with your compiler.

Advice on writing: Whenever you write, be thoughtful about your expected audience, and most importantly *choose at least one specific audience you are writing for*.⁵ For this assignment, imagine that Diamondback is a custom language that your team maintains. You are writing this design document to help with the onboarding process for people joining your team. They are programmers who know generally how compilers work and how x86-64 works, but not how Diamondback has made its choices internally. They will be curious how things work, why decisions were made the way they were, and what upcoming work on the compiler might look like.

Rubric

Note that half of the credit on this assignment is for testing and writing, and the other half is for automatically-graded behavior:

- Required tests: 20% (5 each)
- Behavior of your compiler on our tests: 50%
- Design doc on calling convention: 30%

⁵Writing for multiple audiences at the same time is of course possible, hence “at least” one. It just requires more care.

Extensions

These are not for credit, but you may enjoy thinking through, implementing, or discussing in office hours or with your peers after the assignment is complete.

1. Support tail calls.
2. In some cases, return type annotations can be omitted to reduce work on the programmer. Make the return type annotations optional on function definitions, and calculate them as part of building the definitions environment. Report an error if the type cannot be calculated without the annotation.