

# UCSD CSE131 F19 – Garter

November 21, 2019

**Checkpoint Due Date:** 11pm Wednesday, November 27

**Final Due Date:** 11pm **Thursday** December 5

The specific features listed for the checkpoint are **Open to Collaboration** (detailed below), and the rest is **Closed to Collaboration**.

You will implement memory management atop a type-checked language with heap-allocated data and functions.

Classroom: FILL      Github: FILL

## Syntax

The concrete syntax and type language for Garter is below. We use  $\dots$  to indicate *zero or more* of the previous element. There are boxes around the new pieces of concrete syntax.

$e$	$:=$	$n \mid \text{true} \mid \text{false} \mid x$ $\mid (\text{let } ((x \ e) \ (x \ e) \ \dots) \ e \ e \dots)$ $\mid (\text{if } e \ e \ e)$ $\mid (op_2 \ e \ e) \mid (op_1 \ e)$ $\mid (\text{while } e \ e \ e \dots) \mid (\text{set } x \ e)$ $\mid (f \ e \dots) \mid \boxed{(\text{null } \tau)}$ $\mid \boxed{(\text{get } e \ n)} \mid \boxed{(\text{update } e \ n \ e)}$	$\tau$	$:=$	$\text{Num} \mid \text{Bool} \mid C$
			$\delta$	$:=$	$\text{fun} \mid \text{data}$
			$\Delta$	$:=$	$\{\delta \ f : \tau \dots \rightarrow \tau, \dots\}$
			$\Delta[f]$	means	look up the type of $f$ in $\Delta$
			$\Gamma$	$:=$	$\{x : \tau, \dots\}$
			$\Gamma[x]$	means	look up the type of $x$ in $\Gamma$
$d$	$:=$	$(\text{def } f \ (x : \tau \dots) : \tau \ e \ e \dots)$ $\mid \boxed{(\text{data } C \ (\tau \dots))}$	$(x, \tau) :: \Gamma$	means	add $x$ to $\Gamma$ with type $\tau$
$p$	$:=$	$d \dots e$	$\Delta; \Gamma \vdash e : \tau$	means	with definitions $\Delta$ and env $\Gamma$ , $e$ has type $\tau$
$op_1$	$:=$	$\text{add1} \mid \text{sub1} \mid \text{isNum} \mid \text{isBool} \mid \text{print}$	$\Delta \vdash_d d : \checkmark$	means	with definitions $\Delta$ the definition $d$ type-checks
$op_2$	$:=$	$+\mid-\mid*\mid<\mid>\mid==\mid\boxed{=}$	$\vdash_p p : \checkmark$	means	the program $p$ type checks
$n$	$:=$	63-bit signed number literals			
$x, f, C$	$:=$	variable, function, and constructor names			

## Semantics

The semantics here are all provided for you, we describe them so you'll be able to write accurate tests.

## Data Definitions, Construction, and Manipulation

The main new feature in Garter is data definitions  $(\text{data } C \ (\tau \dots))$ , where  $s$  is the name of the data definition and the types  $\tau \dots$  are the types of the elements stored in instances of the data definition. Elements are accessed and updated positionally

with *fixed* (not computed, as with arrays) numeric indices using `(get e n)` and `(update e n e)`.<sup>1</sup> The syntax for function applications is used to construct new data instances.

As an example, this program evaluates to 67:

```
(data Pair (Num Num))
(let (
  (p1 (Pair 4 5))
  (p2 (Pair 4 5))
  (p3 p1)
)
(update p1 0 11)
(update p2 1 56)
(+ (get p3 0) (get p2 1)))
```

## Printing Data Instances

In Egg-Eater, locations (referring to instances of data) are a new kind of value that can be printed, just like numbers and booleans.

When an instance is printed, it prints in the format

$(C\ v_1\ v_2\ \dots)$

Where  $C$  is the name of the constructor used to create it, and values  $v_1$  and  $v_2$  are the printed form of the values stored in its fields, separated by spaces.

For example:

```
(data Pair (Num Num))
(data PairOfPairs (Pair Pair))
(let ((p (PairOfPairs (Pair (+ 1 2) 6) (Pair (add1 6) 8))))
  p)

# prints:
(PairOfPairs (Pair 3 6) (Pair 7 8))
```

## Equality

There two types of equality in Egg-Eater, reflecting the new nuances of heap-allocated data. The first, `==`, behaves as before on existing values, and on locations referring to instances of data, returns `true` if the *locations* are identical. The second, `=`, behaves as before on existing values, and on locations returns `true` if the two instances came from the same constructor and the *contents* of those locations are all equal according to `=`.

For example:

```
(data Pair (Num Num))
(data PairOfPair (Pair Pair))
(data Point (Num Num))
(let (
  (p1 (Pair 3 4))
  (p2 (Pair 3 4))
  (p3 (Point 3 4))
  (p4 (Point 3 5))
  (pp12 (PairOfPair p1 p2))
```

---

<sup>1</sup>As an analogy, data definitions are somewhat like structs in C, but use positional lookup instead of names; as another analogy, data definitions are like tuples in OCaml and we can match on them by statically known positions using functions like `fst` and `snd`, but not compute the position of lookup.

```

(pp21 (PairOfPair p2 p1))
)
(print (= p1 p2)) ; true, same constructor and contents
(print (== p1 p2)) ; false, different locations
(print (= p1 p3)) ; false, different constructors
(print (== p1 p3)) ; false, different locations
(print (= p3 p4)) ; false, different contents
(print (= pp12 pp21)) ; true, same (nested) contents
0
)

```

## Type Checking

Garter has essentially the same type rules as Diamondback for expressions. The definitions environment  $\Delta$  is constructed with the types of the constructors for data definitions as well as function definitions; these are distinguished by either **data** or **fun** before the name. As an example, the definition (**data** **Point** (**Num** **Num**)) would appear in  $\Delta$  as **data** **Point** : (**Num** **Num**  $\rightarrow$  **Point**). There is a new rule for each new syntactic form, except for **data** definitions which don't need separate type checking.

$$\begin{array}{c}
\text{TR-NULL} \quad \frac{\Delta[\text{data } C] = (\tau_1 \cdots \rightarrow \tau_r)}{\Delta; \Gamma \vdash (\text{null } C) : C} \\
\\
\text{TR-GET} \quad \frac{\Delta; \Gamma \vdash e : C \quad \Delta[\text{data } C] = (\tau_1 \cdots \tau_n \tau_{n+1} \cdots \rightarrow \tau_r)}{\Delta; \Gamma \vdash (\text{get } e \ n) : \tau_n} \\
\\
\text{TR-UPDATE} \quad \frac{\Delta; \Gamma \vdash e : C \quad \Delta[\text{data } C] = (\tau_1 \cdots \tau_n \tau_{n+1} \cdots \rightarrow \tau_r) \quad \Delta; \Gamma \vdash e_v : \tau_n}{\Delta; \Gamma \vdash (\text{update } e \ n \ e_v) : \tau_n}
\end{array}$$

There are a few important features here.

- The **null** expression comes with a type that it should be treated as. The type checker simply checks that this annotation is some **data** type and treats the **null** value as that type. This allows us to construct instances of recursively-defined datatypes like (**Link** (**Num** **Link**)).
- In TR-Get and TR-Update, we check that the first expression has a type of some data definition  $C$ . The types before the  $\rightarrow$  are the types of the fields or elements listed in the data definition.
- We assume the existing rule for TR-App in applications, which simply checks that values with the right types are present in order according to the data definition (just like for function calls).

## Application Binary Interface

### Value and Heap Layout

The value layout is extended to keep track of information needed in garbage collection:

- 0xFFFFFFFFFFFFFFFF[xx1] - Number
- 0x0000000000000000[0110] - True
- 0x0000000000000000[0010] - False
- 0x0000000000000000[0000] - Null

- `0xFFFFFFFFXXXX[x000]` - Data Reference, an address of a data instance on the heap laid out as follows (each set of `[]` is one 8-byte word)

```
[ GC word ][ name reference ][ element count n ][ value 1 ][ value 2 ] ... [ value n ]
```

The use of the GC word is completely up to your memory management implementation and is always initialized to 0 (see below). The name reference is the address of a C string that holds the struct's name (essentially a `char*`) used in printing and equality. The element count tracks the number of elements stored in the data value.

As an example, consider this program:

```
(data Pair (Num Num))
(let (
  (p1 (Pair 4 5))
  (p2 (Pair 6 7))
  (p3 p1)
  )
  ...)
```

The stack word for `p1` would hold a value like `0x00000000ABCDE120`, where at address `0x00000000ABCDE120` would be stored:

```
0x00000000ABCDE230 : [ 0x0000000000000000 ] ; gc word
                    [ 0x00000000NAMEADDR ] ; address of "Pair"
                    [ 0x0000000000000002 ] ; count of elements
                    [ 0x0000000000000009 ] ; representation of 4
                    [ 0x000000000000000B ] ; representation of 5
```

Where at `0xNAMEADDR` we would find the characters `Pair\0`, and 9 and 11 are the representations of 4 and 5. At the stack word for `p3` we would also find `0x00000000ABCDE120`. At the stack word for `p2` we should expect to find a different address, say `0x00000000ABCDE230`, with a similar layout but different values:

```
0x00000000ABCDE230 : [ 0x0000000000000000 ] ; gc word
                    [ 0x00000000NAMEADDR ] ; address of "Pair"
                    [ 0x0000000000000002 ] ; count of elements
                    [ 0x000000000000000D ] ; representation of 6
                    [ 0x000000000000000F ] ; representation of 7
```

## Calling Convention

We use a calling convention similar to the one discussed in class, so at any given moment there are a number of function calls on the stack, each with arguments and local variables.

Some important highlights:

- On the right, we show the addresses stored in the arguments given to `try_gc` which are passed on to the `gc` function you will write. This includes `stack_top`, which is equal to `rsp - (stackloc si)`, `first_frame`, which is equal to `rsp`, and `STACK_BOTTOM`, which is a global that refers to the original value of `rsp` right after calling `our_code_starts_here`. We will say more about each of these in the next section.
- We made sure the compiler implements the invariant that `rsp - (stackloc si)` will always refer to the word above the topmost valid value, and that there won't be any invalid values in the local variables or the arguments on the stack.

## The Garbage Collection Interface

The generated code from Garter checks, on each allocation of a new `data` value, if enough space is left on the heap to fit that value. It does this by comparing the value in `r15` to the value stored in `HEAP_END`, which is a global variable set by `main`. If

```

[   UNUSED SPACE   ] <- stack_top
-----
[local var N       ]
[...               ]
[local var 1       ]   these locals and args are
[arg 1             ]   for the topmost active
[...]             ]   function call
[arg N             ]
[prev rsp value    ]
rsp -> [return address] <- first_frame
-----
...
-----
[local var N       ]   these locals and args are
[...]             ]   for a current active
[local var 1       ]   function call
[arg 1             ]
[...]             ]
[arg N             ]
[prev rsp value    ]
[return address    ]
-----
[local var N       ]   these locals are for
[...]             ]   the main expression
[local var 1       ]
[ret ptr to main   ] <- STACK_BOTTOM

```

enough space isn't available, the generated code calls the `try_gc` function with several arguments:

- `alloc_ptr` – the current value in `r15`
- `words_needed` – the number of words needed by the allocation
- `first_frame` – the address of the stored `rsp` in the topmost (most-recently-called) stack frame of a Garter function (also the most recent value of `rsp` used by Garter)
- `stack_top` – the address of the word immediately above the stack used by Garter; also equal to `rsp - (stackloc si)` at the time of allocation.

The `try_gc` function uses this information, along with several global variables, to call the `gc` function, which you will implement. The `gc` function receives

- `STACK_BOTTOM`, which is the address of the bottom-most word on the Garter stack, which corresponds to the return pointer to `main` used at the end of `our_code_starts_here`
- `first_frame` and `stack_top` as provided to `try_gc`
- The current value stored in `HEAP`. `HEAP` is initially set to the address allocated for the start of the heap in `main` on program startup. Your program can (but does not need to) manipulate this.
- The current value stored in `HEAP_END`. `HEAP_END` is initially set to the address immediately *after* the end of the heap in `main` on program startup. Your program can (but does not need to) manipulate this.

The `gc` function is expected to return a new value for the heap pointer.